

Online Optimization

Lecturer: Jiaming Liang

November 14, 2024

1 Online linear optimization

The online linear optimization problem can be described as the following game. We are given a closed convex set $S \subseteq \mathbb{R}^n$ (strategy space) and we select some strategy $x_t \in S$ in each round $t \geq 0$. We are then revealed (by the environment or the adversary) a penalty $p_t \in \mathbb{R}^n$. Our goal is to minimize the regret of T rounds of this game:

$$\text{regret}(T) = \sum_{t=0}^{T-1} p_t^\top x_t - \min_{x \in S} \sum_{t=0}^{T-1} p_t^\top x. \quad (1)$$

The online convex optimization problem generalizes online linear optimization in that after choosing a decision $x_t \in S$, the penalty is $f_t(x_t)$ instead of $p_t^\top x_t$, where f_t is convex. The regret of T rounds becomes

$$\text{regret}(T) = \sum_{t=0}^{T-1} f_t(x_t) - \min_{x \in S} \sum_{t=0}^{T-1} f_t(x).$$

We focus the online linear optimization problem in the rest of the lecture. The ultimate goal is to design algorithms so that $\lim_{T \rightarrow \infty} \text{regret}(T)/T = 0$, i.e., the average regret over the rounds converges to 0. Such an algorithm is said to have sublinear regret as $\text{regret}(T) = o(T)$.

Lemma 1. Consider a convex function $f : S \rightarrow \mathbb{R}$ and assume that f has a subgradient oracle and minimizer $x_* = \operatorname{argmin}_{x \in S} f(x)$. If $p_t \in \partial f(x_t)$ for every $t \geq 0$, letting $f_* = f(x_*)$, then we have

$$f\left(\frac{1}{T} \sum_{t=0}^{T-1} x_t\right) - f_* \leq \frac{\text{regret}(T)}{T}.$$

Proof. It follows from the definition of regret in (1) and the convexity of f that

$$\begin{aligned} f\left(\frac{1}{T} \sum_{t=0}^{T-1} x_t\right) - f_* &\leq \frac{1}{T} \sum_{t=0}^{T-1} f(x_t) - f_* \leq \frac{1}{T} \sum_{t=0}^{T-1} p_t^\top (x_t - x_*) \\ &\leq \frac{1}{T} \left[\sum_{t=0}^{T-1} p_t^\top x_t - \min_{x \in S} \sum_{t=0}^{T-1} p_t^\top x \right] = \frac{\text{regret}(T)}{T}. \end{aligned}$$

□

We note that the above lemma gives an important reduction: if we have an algorithm to solve online linear optimization, given $p_t \in \partial f(x_t)$ for every $t \geq 0$, then the same algorithm also solves the offline optimization $\min_{x \in S} f(x)$.

2 Follow the regularized leader

In online linear optimization, a standard algorithm is the follow the regularized leader (FTRL): pick a regularizer $r : \mathbb{R}^n \rightarrow \mathbb{R}$, that is differentiable and μ -strongly convex on S with respect to some norm $\|\cdot\|$, then FTRL generates x_T as follows.

Algorithm 1 FTRL

Input: pick an initial strategy $x_0 \in S$

for $T \geq 1$ **do**

 Step 1. Environment reveals a penalty $p_{T-1} \in \mathbb{R}^n$.

 Step 2. Compute $x_T = \operatorname{argmin}_{x \in S} \left\{ \Phi_T(x) := \eta \sum_{t=0}^{T-1} p_t^\top x + r(x) \right\}$.

end for

Note that Φ_T is defined as in FTRL for $T \geq 1$ and $\Phi_0(x) = r(x)$. We assume there exists an oracle to solve the subproblem in step 2 of FTRL. For example, if $r(x)$ is ℓ_2 -norm squared, then the oracle for step 2 is just proj_S . It is interesting to observe that FTRL is the same as the constant stepsize dual averaging (see Lecture 9). As a result, the regret analysis of FTRL is very close to the convergence analysis of dual averaging, while we still present the analysis for completeness.

Lemma 2. *If $\|p_t\|_* \leq G$ for all $t \geq 0$, applying FTRL on the online linear optimization problem, then for every $T \geq 0$, we have*

$$-\frac{\eta^2 G^2 T}{2\mu} + \eta \sum_{t=0}^{T-1} p_t^\top x_t + \min_{x \in S} r(x) \leq \min_{x \in S} \Phi_T(x).$$

Proof. Proof by induction. Since $\Phi_T(x) = r(x)$, the case $T = 0$ is trivial. Assume the claim is true for some $T \geq 0$. Using the definition of Φ_T in FTRL and the induction hypothesis, we have

$$\begin{aligned} \Phi_{T+1}(x_{T+1}) &= \Phi_T(x_{T+1}) + \eta p_T^\top x_{T+1} \\ &\geq \Phi_T(x_T) + \frac{\mu}{2} \|x_{T+1} - x_T\|^2 + \eta p_T^\top x_{T+1} \\ &\geq -\frac{\eta^2 G^2 T}{2\mu} + \eta \sum_{t=0}^{T-1} p_t^\top x_t + \min_{x \in S} r(x) + \eta \left[p_T^\top x_{T+1} + \frac{\mu}{2\eta} \|x_{T+1} - x_T\|^2 \right], \end{aligned}$$

where the first inequality is due to the fact that Φ_T is μ -strongly convex in $\|\cdot\|$. It follows from the Cauchy-Schwarz inequality and the assumption that $\|p_t\|_* \leq G$ that

$$p_T^\top x_{T+1} + \frac{\mu}{2\eta} \|x_{T+1} - x_T\|^2 \geq -\|p_T\|_* \|x_{T+1} - x_T\| + \frac{\mu}{2\eta} \|x_{T+1} - x_T\|^2 + p_T^\top x_T \geq -\frac{\eta G^2}{2\mu} + p_T^\top x_T.$$

Therefore, combining the above two equations, we verify that the claim holds for the case $T+1$. \square

Theorem 1. *If $\|p_t\|_* \leq G$ for every $t \geq 0$, applying FTRL on the online linear optimization problem, then for every $T \geq 1$, we have the following regret bound*

$$\text{regret}(T) \leq \frac{\eta G^2 T}{2\mu} + \frac{1}{\eta} \left(\max_{x \in S} r(x) - \min_{x \in S} r(x) \right).$$

Moreover, if $p_t \in \partial f(x_t)$ for every $t \geq 0$ and $\max_{x \in S} r(x) - \min_{x \in S} r(x) \leq R^2$ for some $R > 0$, and using stepsize $\eta = \sqrt{(2R^2\mu)/(TG^2)}$, then the regret of FTRL is bounded as follows

$$\text{regret}(T) \leq RG \sqrt{\frac{2T}{\mu}},$$

and FTRL gives an optimality gap of $\min_{x \in S} f(x)$ as follows

$$f\left(\frac{1}{T} \sum_{t=0}^{T-1} x_t\right) - f_* \leq RG \sqrt{\frac{2}{T\mu}}.$$

Proof. It follows from Lemma 2 and the definition of Φ_T that for every $x \in S$

$$-\frac{\eta^2 G^2 T}{2\mu} + \eta \sum_{t=0}^{T-1} p_t^\top x_t + \min_{x \in S} r(x) \leq \eta \sum_{t=0}^{T-1} p_t^\top x + r(x).$$

Rearranging the terms, we have

$$\sum_{t=0}^{T-1} p_t^\top (x_t - x) \leq \frac{\eta G^2 T}{2\mu} + \frac{1}{\eta} \left(r(x) - \min_{x \in S} r(x) \right).$$

Maximizing over $x \in S$ and using the definition of regret in (1), we obtain

$$\text{regret}(T) \leq \frac{\eta G^2 T}{2\mu} + \frac{1}{\eta} \left(\max_{x \in S} r(x) - \min_{x \in S} r(x) \right).$$

Using stepsize $\eta = \sqrt{(2R^2\mu)/(TG^2)}$ and bound R^2 , we derive the regret of FTRL bounded by

$$\text{regret}(T) \leq RG \sqrt{\frac{2T}{\mu}}.$$

It immediately follows from the above regret of FTRL and Lemma 1 that

$$f\left(\frac{1}{T} \sum_{t=0}^{T-1} x_t\right) - f_* \leq \frac{\text{regret}(T)}{T} = \frac{RG \sqrt{2T/\mu}}{T} = RG \sqrt{\frac{2}{T\mu}}.$$

\square

Clearly, the complexity for FTRL to find an ε -solution to $\min_{x \in S} f(x)$ is

$$\mathcal{O}\left(\frac{R^2 G^2}{\mu \varepsilon^2}\right),$$

which matches that of dual averaging for solving nonsmooth optimization problems.

As we know dual averaging is also applicable to smooth optimization problems, we now use FTRL to recover the convergence of projected gradient-type methods (such as dual averaging and the projected gradient method) for minimizing smooth functions.

Lemma 3. *Assume that f is L -smooth over S , and let $r(x) = \|x - x_0\|_2^2/2$ in step 2 of FTRL and let $p_t = \nabla f(x_t)$ for every $t \geq 0$. Applying FTRL with $\eta = 1/L$ on the online linear optimization problem, then for every $T \geq 0$, we have*

$$\eta \sum_{t=1}^T f(x_t) - \eta \sum_{t=0}^{T-1} \left(f(x_t) - p_t^\top x_t \right) \leq \min_{x \in S} \Phi_T(x)$$

Proof. Proof by induction. Since $\Phi_0(x) = \|x - x_0\|_2^2/2$, the case $T = 0$ is trivial. Assume the claim is true for some $T \geq 0$. Using the definition of Φ_T in FTRL and the induction hypothesis, we have

$$\begin{aligned} \Phi_{T+1}(x_{T+1}) &= \Phi_T(x_{T+1}) + \eta p_T^\top x_{T+1} \\ &\geq \Phi_T(x_T) + \frac{1}{2} \|x_{T+1} - x_T\|_2^2 + \eta p_T^\top x_{T+1} \\ &\geq \eta \sum_{t=1}^T f(x_t) - \eta \sum_{t=0}^{T-1} \left(f(x_t) - p_t^\top x_t \right) + \eta \left[p_T^\top x_{T+1} + \frac{1}{2\eta} \|x_{T+1} - x_T\|_2^2 \right], \end{aligned}$$

where the first inequality is due to the fact that Φ_T is 1-strongly convex in $\|\cdot\|_2$. It follows from the fact that $p_t = \nabla f(x_t)$ and $\eta = 1/L$, and the assumption that f is L -smooth that

$$p_T^\top x_{T+1} + \frac{1}{2\eta} \|x_{T+1} - x_T\|_2^2 = \nabla f(x_T)^\top x_{T+1} + \frac{L}{2} \|x_{T+1} - x_T\|_2^2 \geq f(x_{T+1}) - \left(f(x_T) - \nabla f(x_T)^\top x_T \right).$$

Therefore, combining the above two equations, we verify that the claim for the case $T+1$ holds. \square

Theorem 2. *Assuming the conditions in Lemma 3 hold and S has a diameter $D > 0$, and applying FTRL with $\eta = 1/L$ on the online linear optimization problem, then for every $T \geq 1$, we have the following regret bound*

$$\text{regret}(T) \leq f(x_0) - f_* + \frac{LD^2}{2},$$

and FTRL gives an optimality gap of $\min_{x \in S} f(x)$ as follows

$$f\left(\frac{1}{T} \sum_{t=0}^{T-1} x_t\right) - f_* \leq \frac{2[f(x_0) - f_*] + LD^2}{2T}.$$

Proof. It follows from Lemma 3 that for every $x \in S$

$$\eta[f(x_T) - f(x_0)] + \eta \sum_{t=0}^{T-1} p_t^\top x_t \leq \eta \sum_{t=0}^{T-1} p_t^\top x + \frac{1}{2} \|x - x_0\|_2^2.$$

Rearranging the terms and using the fact that $\eta = 1/L$ yields

$$\sum_{t=0}^{T-1} p_t^\top (x_t - x) \leq f(x_0) - f(x_T) + \frac{L}{2} \|x - x_0\|_2^2.$$

Maximizing over $x \in S$ and using the definition of regret in (1) and the boundedness of S , we obtain

$$\text{regret}(T) \leq f(x_0) - f_* + \frac{LD^2}{2}.$$

It immediately follows from the above regret of FTRL and Lemma 1 that

$$f\left(\frac{1}{T} \sum_{t=0}^{T-1} x_t\right) - f_* \leq \frac{\text{regret}(T)}{T} \leq \frac{2[f(x_0) - f_*] + LD^2}{2T}.$$

□

The following theorem gives a different convergence result of FTRL without using the reduction, i.e., Lemma 1. It recovers the convergence rate of dual averaging for solving smooth optimization problems.

Theorem 3. *Assuming the conditions in Lemma 3 hold and applying FTRL with $\eta = 1/L$ on the online linear optimization problem, then for every $T \geq 1$, we have*

$$f\left(\frac{1}{T} \sum_{t=1}^T x_t\right) - f_* \leq \frac{L\|x_0 - x_*\|_2^2}{2T}.$$

Proof. It immediately follows from the fact that $p_t = \nabla f(x_t)$ for every $t \geq 0$ and Lemma 3 that for every $x \in S$

$$\eta \sum_{t=1}^T f(x_t) - \eta \sum_{t=0}^{T-1} \left(f(x_t) - \nabla f(x_t)^\top x_t \right) \leq \eta \sum_{t=0}^{T-1} \nabla f(x_t)^\top x + \frac{1}{2} \|x - x_0\|_2^2.$$

Rearranging the terms and using the fact that $\eta = 1/L$ yields

$$\sum_{t=1}^T f(x_t) \leq \sum_{t=0}^{T-1} (f(x_t) + \langle \nabla f(x_t), x - x_t \rangle) + \frac{L}{2} \|x - x_0\|_2^2 \leq Tf(x) + \frac{L}{2} \|x - x_0\|_2^2,$$

where the second inequality is due to the convexity of f . Taking $x = x_*$ in the above inequality and using the convexity of f again, we obtain

$$f\left(\frac{1}{T}\sum_{t=1}^T x_t\right) - f_* \leq \frac{L\|x_0 - x_*\|_2^2}{2T}.$$

□

Note that Theorem 3 shows convergence at $\frac{1}{T}\sum_{t=1}^T x_t$, which is different from $\frac{1}{T}\sum_{t=0}^{T-1} x_t$ in Theorems 1 and 2.

3 Saddle point problem

In Section 2, we have seen FTRL originally designed for online linear optimization can be readily applied to offline optimization, both smooth and nonsmooth minimization problems. We discuss another interesting application of FTRL in this section, that is, the saddle point problem.

The saddle point problem considered in this section is as follows

$$\min_{x \in \mathcal{X}} \max_{y \in \mathcal{Y}} f(x, y)$$

where $\mathcal{X} \subseteq \mathbb{R}^n$ and $\mathcal{Y} \subseteq \mathbb{R}^m$ are closed convex sets and $f : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$ is differentiable and convex-concave (i.e., $f(\cdot, y)$ is convex for any $y \in \mathcal{Y}$ and $f(x, \cdot)$ is concave for any $x \in \mathcal{X}$). It is equivalent to

$$\min_{x \in \mathcal{X}} \left\{ f_{\mathcal{X}}(x) := \max_{y \in \mathcal{Y}} f(x, y) \right\}.$$

For simplicity, we denote $f(z) = f(z_{\mathcal{X}}, z_{\mathcal{Y}})$, where $z \in \mathcal{Z} = \mathcal{X} \times \mathcal{Y}$ and $z_{\mathcal{X}}$ (resp., $z_{\mathcal{Y}}$) denotes the \mathcal{X} -component (resp., \mathcal{Y} -component) of z .

For every $z \in \mathcal{Z}$, we define the duality gap as

$$\text{gap}(z) = \text{gap}(z_{\mathcal{X}}, z_{\mathcal{Y}}) = \max_{y \in \mathcal{Y}} f(z_{\mathcal{X}}, y) - \min_{x \in \mathcal{X}} f(x, z_{\mathcal{Y}}) = f_{\mathcal{X}}(z_{\mathcal{X}}) - f_{\mathcal{Y}}(z_{\mathcal{Y}}),$$

where $f_{\mathcal{Y}}(y) := \min_{x \in \mathcal{X}} f(x, y)$. We say that $z \in \mathcal{Z}$ is an ε -Nash equilibrium, or has ε -duality gap, or is ε -optimal if $\text{gap}(z) \leq \varepsilon$. Further, we call $z \in \mathcal{Z}$ a Nash equilibrium if $\text{gap}(z) = 0$. An intuitive interpretation of the duality gap is as follows: consider there are two players x and y in a zero-sum game, the goal of x is to minimize $f(x, y)$ and hence save cost, while the goal of y is to maximize $f(x, y)$ and hence make profit. The duality gap can be written as

$$\text{gap}(z) = \left[\max_{y \in \mathcal{Y}} f(z_{\mathcal{X}}, y) - f(z) \right] + \left[f(z) - \min_{x \in \mathcal{X}} f(x, z_{\mathcal{Y}}) \right],$$

where the first gap term is the profit that y makes with $z_{\mathcal{X}}$ being fixed and the second gap term is the cost that x saves with $z_{\mathcal{Y}}$ being fixed. A Nash equilibrium is a pair of strategies z at which

both x and y are not willing to change their individual strategies $z_{\mathcal{X}}$ and $z_{\mathcal{Y}}$, as none of them would increase their own utility functions (i.e., saving cost and making profit) by doing so.

Note that $f_{\mathcal{X}}(z_{\mathcal{X}}) \geq f(z) \geq f_{\mathcal{Y}}(z_{\mathcal{Y}})$ for every $z \in \mathcal{Z}$ by definition, and hence $f_{\mathcal{X}}^* \geq f_{\mathcal{Y}}^*$ where $f_{\mathcal{X}}^* = \min_{x \in \mathcal{X}} f_{\mathcal{X}}(x)$ and $f_{\mathcal{Y}}^* = \max_{y \in \mathcal{Y}} f_{\mathcal{Y}}(y)$. This is indeed the weak duality and it is equivalent to

$$\min_{x \in \mathcal{X}} \max_{y \in \mathcal{Y}} f(x, y) \geq \max_{y \in \mathcal{Y}} \min_{x \in \mathcal{X}} f(x, y).$$

If $f(x, y)$ is convex-concave, then “=” holds. Clearly, $\text{gap}(z) \geq f_{\mathcal{X}}(z_{\mathcal{X}}) - f_{\mathcal{X}}^* + f_{\mathcal{Y}}^* - f_{\mathcal{Y}}(z_{\mathcal{Y}})$ and therefore if $z \in \mathcal{Z}$ is an ε -Nash equilibrium, then $z_{\mathcal{X}}$ (resp., $z_{\mathcal{Y}}$) is an ε -solution for $\min_{x \in \mathcal{X}} f_{\mathcal{X}}$ (resp., $\max_{y \in \mathcal{Y}} f_{\mathcal{Y}}$). Furthermore, this reasoning implies that if there exists a Nash equilibrium, i.e., $z^* \in \mathcal{Z}$ with $\text{gap}(z^*) = 0$, then $f_{\mathcal{X}}(z_{\mathcal{X}}^*) = f_{\mathcal{X}}^*$ and $f_{\mathcal{Y}}(z_{\mathcal{Y}}^*) = f_{\mathcal{Y}}^*$ and correspondingly $f_{\mathcal{X}}^* = f_{\mathcal{Y}}^*$. This is called the strong duality.

The discussions above reflect some key ideas in Lecture 9.

Analogous to the reduction from online linear optimization to offline optimization (see Lemma 1), we next develop the reduction from online linear optimization to the saddle point problem. Thus, we can apply FTRL to solve the saddle point problem.

Lemma 4. Let $g : \mathcal{X} \times \mathcal{Y} \rightarrow \mathcal{X} \times \mathcal{Y}$ denote the gradient for $f(x, y)$, i.e., $g(x, y) = (\nabla_{\mathcal{X}} f(x, y), -\nabla_{\mathcal{Y}} f(x, y))$. For $z^0, \dots, z^{T-1} \in \mathcal{Z}$ and $\bar{z} = \frac{1}{T} \sum_{t=0}^{T-1} z^t$, we have

$$\text{gap}(\bar{z}) \leq \text{regret}(T).$$

Proof. Since f is convex-concave, we have for every $u \in \mathcal{Z}$,

$$\begin{aligned} f(u_{\mathcal{X}}, z_{\mathcal{Y}}^t) &\geq f(z^t) + \nabla_{\mathcal{X}} f(z^t)^{\top} (u_{\mathcal{X}} - z_{\mathcal{X}}^t), \\ f(z_{\mathcal{X}}^t, u_{\mathcal{Y}}) &\leq f(z^t) + \nabla_{\mathcal{Y}} f(z^t)^{\top} (u_{\mathcal{Y}} - z_{\mathcal{Y}}^t). \end{aligned}$$

Combining the above two inequalities yields that

$$g(z^t)^{\top} (z^t - u) \geq f(z_{\mathcal{X}}^t, u_{\mathcal{Y}}) - f(u_{\mathcal{X}}, z_{\mathcal{Y}}^t).$$

Summing the above inequality from $t = 0$ to $T - 1$ and applying Jensen's inequality, we have

$$\frac{1}{T} \sum_{t=0}^{T-1} g(z^t)^{\top} (z^t - u) \geq \frac{1}{T} \sum_{t=0}^{T-1} [f(z_{\mathcal{X}}^t, u_{\mathcal{Y}}) - f(u_{\mathcal{X}}, z_{\mathcal{Y}}^t)] \geq f(\bar{z}_{\mathcal{X}}, u_{\mathcal{Y}}) - f(u_{\mathcal{X}}, \bar{z}_{\mathcal{Y}}).$$

Maximizing over $u \in \mathcal{Z}$ and using the definition of regret in (1), we obtain

$$\text{regret}(T) \geq \max_{u \in \mathcal{Z}} f(\bar{z}_{\mathcal{X}}, u_{\mathcal{Y}}) - f(u_{\mathcal{X}}, \bar{z}_{\mathcal{Y}}).$$

It follows from the definition of the duality gap that

$$\text{gap}(\bar{z}) = \max_{y \in \mathcal{Y}} f(\bar{z}_{\mathcal{X}}, y) - \min_{x \in \mathcal{X}} f(x, \bar{z}_{\mathcal{Y}}) = \max_{u \in \mathcal{Z}} f(\bar{z}_{\mathcal{X}}, u_{\mathcal{Y}}) - f(u_{\mathcal{X}}, \bar{z}_{\mathcal{Y}}).$$

The conclusion of the lemma immediately follows. \square

We skip the regret of FTRL and its implication in the saddle point problem for shortness.